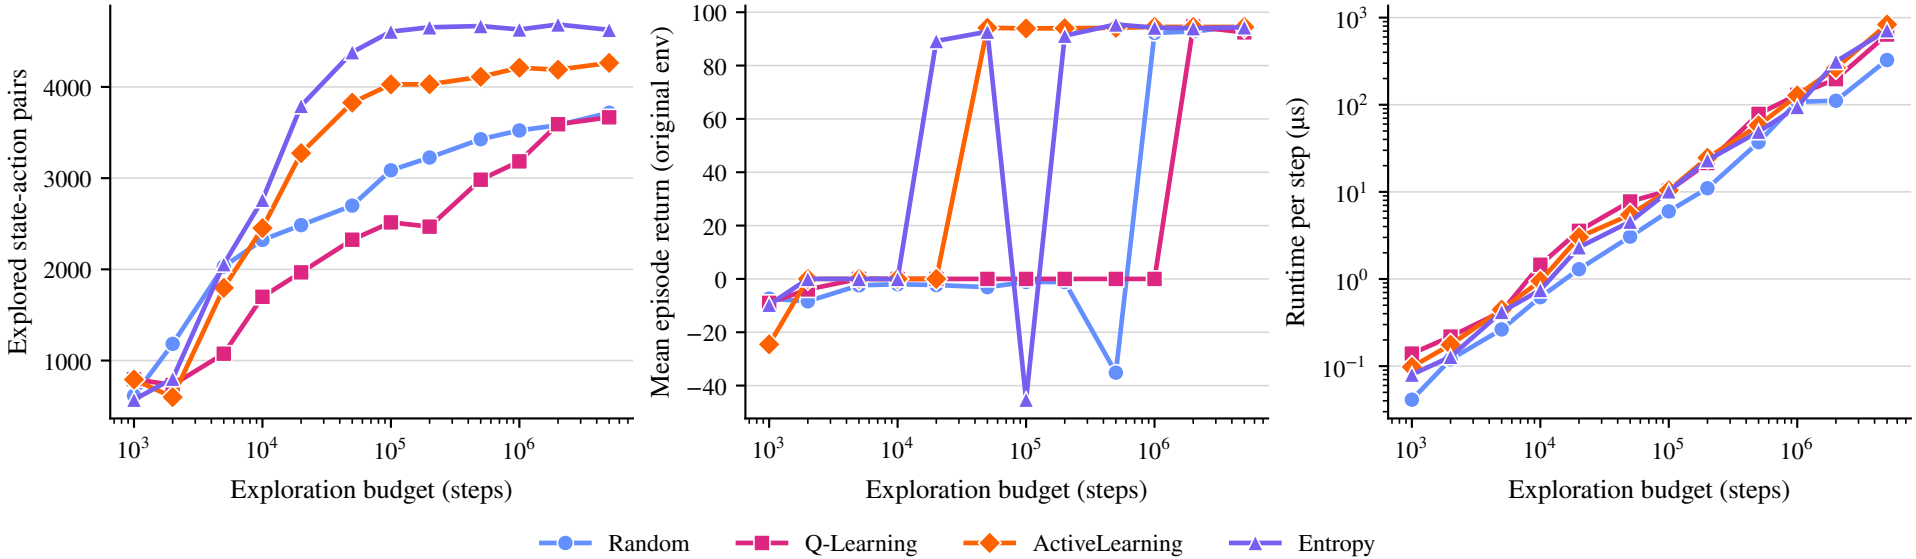
Explored state-action pairs
1000
2000
3000
4000
10^3
10^4
10^5
10^6
Exploration budget (steps)
Mean episode return (original env)
100
80
60
40
20
0
−20
−40
10^3
10^4
10^5
10^6
Exploration budget (steps)
Runtime per step (µs)
10^3
10^2
10^1
10^0
10^{-1}
10^3
10^4
10^5
10^6
Exploration budget (steps)
Random
Q-Learning
ActiveLearning
Entropy